

Efficient Methods for Multi-Qubit Circuit Optimization

Luke Heyfron, and Earl T. Campbell

October 2017

Abstract

1 Introduction

2 Workflow Overview

The high level work-flow of the protocol begins by converting an n qubit input Clifford + T circuit, C_{in} , (see figure 1.1) to a CNOT + T circuit plus additional operations that are all Clifford (figure 1.2). The purpose of this is to isolate the non-Clifford behaviour within a subcircuit that has properties required for our optimization framework (see section 3). The main obstacle to this conversion process is the presence of Hadamard gates in the original circuit. External Hadamards, found at the beginning and end of the circuit, can be left in place. In order to remove internal Hadamards we use an approach similar to that of [1, 2] where a unitary Hadamard gate is replaced with a Hadamard gadget that involves only CNOT + T gates as well as preparation of ancillas in the $|+\rangle$ state and measurement in the Pauli- X basis (see figure 2). In doing so we are trading internal Hadamards for external Hadamards and space-like overhead.

Once the Hadamards are removed, we decompose the CNOT + T circuit into two partitions such that one part consists of CNOT gates only and the other is diagonal in the computational basis (fig. 1.3). We then extract the *signature tensor* of the diagonal CNOT + T circuit, which is an object that uniquely characterizes the unitary U_f up to a Clifford factor.

The core layer of the algorithm, *T-optimiser*, takes as input the signature tensor of the diagonal CNOT + T circuit and outputs a *gate synthesis matrix*, A , which represents a particular circuit decomposition for the input unitary. The number of columns of A is equal to the T count of its corresponding circuit, so *T-optimiser* ideally outputs a gate synthesis matrix with minimal column number. It is believed that finding optimal solutions for this problem is NP-hard [3], so we must use efficient heuristics that yield near-minimal column number.

The output A matrix is mapped to a circuit of the same form as U_f , with which the original U_f circuit is replaced. The result is an output circuit that implements the full Clifford + T input circuit on the first n qubits and has near optimal T count. The following sections detail these steps in turn. New algorithms for the *T-Optimize* subroutine are the main technical contribution of this paper and so are the main focus.

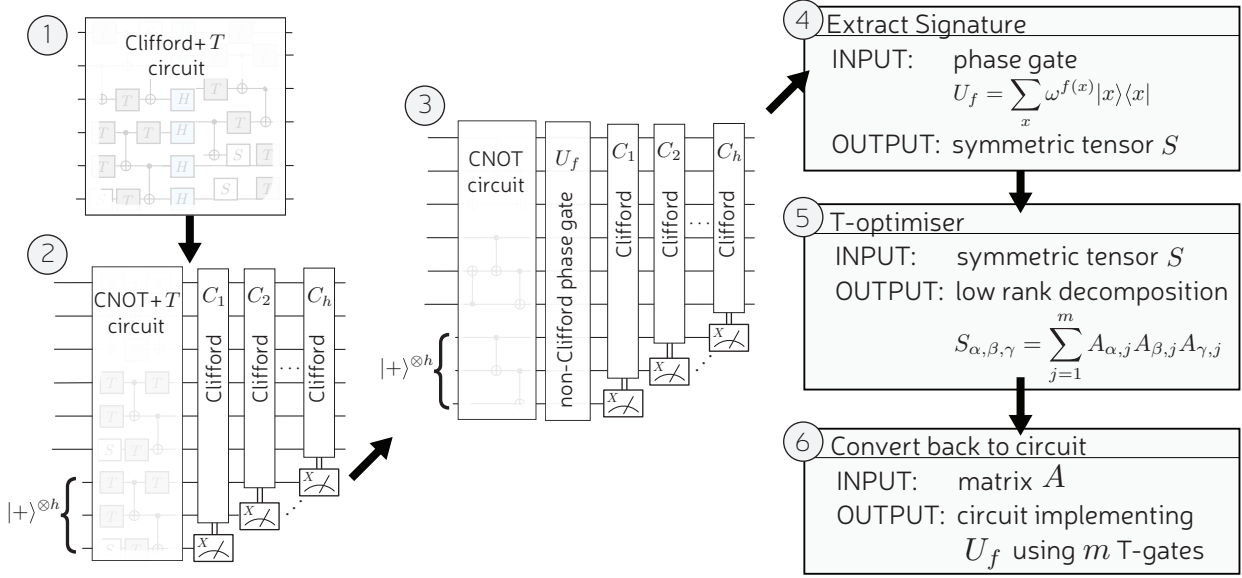


Figure 1: The high level workflow of the T gate optimization protocol is shown. A Clifford + T circuit is converted to the CNOT+T gate set by introducing ancillas and performing classically controlled Clifford gates. A non-Clifford phase gate is extracted, which maps to a signature tensor upon which the core optimization algorithm is performed. The optimized symmetric tensor decomposition is then converted back into a circuit of the form in panel 2) yielding an implementation of the original Clifford + T circuit with reduced T count.




Figure 2: Internal Hadamards are removed by replacing each one according to this circuit identity. The Hadamards on the left-hand circuit are now external and act on an ancilla, therefore they no longer obstruct the *T-optimiser* algorithms. The classically controlled Pauli-X correction of each Hadamard gadget must be conjugated through to the end of the circuit in order to maximize the extent of the CNOT+T sub-circuit of figure 1.2. Because circuits generated by CNOT+T form a group within the third level of the Clifford hierarchy, the conjugation of a Pauli operator through this circuit leads to classically controlled Clifford operators C_1, C_2, \dots, C_h in figure 1.2 and 1.3.

3 Diagonal CNOT+T Framework

This section describes how we map the gate synthesis problem from the quantum circuit picture to an abstract mathematical picture involving binary tensors. We proceed by borrowing the observation from [4],[5] that the action of any n -qubit diagonal CNOT + T unitary U_f on a computational basis state $|\mathbf{x}\rangle$ is given by

$$U_f |\mathbf{x}\rangle = \omega^{f(\mathbf{x})} |\mathbf{x}\rangle, \quad (1)$$

where $\omega = e^{i\frac{\pi}{4}}$, and U_f is completely characterized by a polynomial $f : \mathbb{Z}_2^n \mapsto \mathbb{Z}_8$, known as the phase function of U_f . A phase function can be decomposed into a sum of linear, quadratic and cubic monomials on the x_i where each monomial of order r has a coefficient in \mathbb{Z}_8 and is weighted by a factor 2^{r-1} , as in the following:

$$f(\mathbf{x}) = \sum_{\alpha=1}^n l_{\alpha} x_{\alpha} + 2 \sum_{\alpha < \beta} q_{\alpha,\beta} x_{\alpha} x_{\beta} + 4 \sum_{\alpha < \beta < \gamma} c_{\alpha,\beta,\gamma} x_{\alpha} x_{\beta} x_{\gamma} \pmod{8}, \quad (2)$$

where $l_{\alpha}, q_{\alpha,\beta}, c_{\alpha,\beta,\gamma} \in \mathbb{Z}_8$. In [6] such decompositions of f are referred to as *weighted polynomials* and it was proven that any two unitaries with weighted polynomials whose coefficients all have the same parity are Clifford equivalent. Note the weighted polynomial can be lifted directly from the circuit definition of U_f if we work in the $\{T, CS, CCZ\}$ basis, as each kind of gate corresponds to the linear, quadratic and cubic terms, respectively. We define the *signature tensor*, $S \in \mathbb{Z}_2^{(n,n,n)}$, of U_f to be a symmetric tensor of order 3 whose elements are equal to the

parity of the weighted polynomial coefficients of U_f according to the following relations:

$$S_{\sigma(\alpha,\alpha,\alpha)} = S_{a,a,a} = l_\alpha \pmod{2} \quad (3)$$

$$S_{\sigma(\alpha,\beta,\beta)} = S_{\sigma(\alpha,\alpha,\beta)} = q_{\alpha,\beta} \pmod{2} \quad (4)$$

$$S_{\sigma(\alpha,\beta,\gamma)} = c_{\alpha,\beta,\gamma} \pmod{2} \quad (5)$$

for all $\sigma \in \mathcal{S}_3$, the symmetric group on a set with 3 elements. It follows from this definition that any two unitaries with the same signature tensor are Clifford equivalent.

We recall the definition of gate synthesis matrices from [6], where a matrix, A in $\mathbb{Z}_2^{(n,m)}$, is a gate synthesis matrix for a unitary U_f if it satisfies,

$$f(\mathbf{x}) = |A^T \mathbf{x}| \pmod{8} \quad (6)$$

where f is the phase function of U_f and $|\cdot|$ is the Hamming weight of a binary vector. An A matrix can be efficiently extracted from a diagonal CNOT + T circuit by tracking the action of each gate on the computational basis states through the circuit. Conversely, given an A matrix one can generate a circuit that implements U_f using m T gates. We can determine the signature tensor of U_f from an A matrix of U_f using the following relation,

$$S_{\alpha,\beta,\gamma} = \sum_{j=1}^m A_{\alpha,j} A_{\beta,j} A_{\gamma,j} \pmod{2} \quad (7)$$

We now have the necessary tools to define the gate synthesis problem in the binary tensor picture.

Problem 3.1. Symmetric 3-Tensor Rank (3-STR). Given a unitary U_f whose signature tensor is S , find an A matrix with minimal number of columns that satisfies equation 7.

4 T -optimiser

Until now the T -optimiser subroutine of our protocol has been treated as a black box whose input is a signature tensor S and the output is a gate synthesis matrix S with few columns. In this section we will describe the inner workings of the various T -optimisers we have implemented in this work. It has been shown that the 3-STR problem can be mapped in polynomial time to the minimum distance decoding problem of the punctured Reed-Muller code of order $n-4$ and length n , compactly written as $RM^*(n-4, n)$, and is believed to be a hard problem [3]. This imposes a practical upper bound on the number of qubits n_{RM} over which circuits can be optimally synthesized with respect to T count.

4.1 Reed-Muller decoder (RM)

Although Reed-Muller decoding is believed to be hard, a brute force solver can be implemented for a small number of qubits. We implement such a brute force decoder and found its limit to be 6 qubits. To gain some intuition for the complexity of the problem, consider the following. The number of codespace generators for $RM^*(n-4, n)$ is equal to $N_G = \sum_{r=1}^{n-4} \frac{n!}{[(n-r)!r!]}$. Therefore, the size of the search space is $N_{\text{search}} = 2^{N_G}$. On a processor with a clock speed of 3.20GHz, assuming we can check one codeword per clock cycle, it would take over 91 years to exhaustively search this space and therefore determine optimality for a general CNOT+ T circuit for $n=7$. Performing the same back-of-the-envelope calculation for $n=6$, it would take $\approx 7 \times 10^{-4}$ seconds when in actuality it takes around 2 minutes, meaning the time for $n=7$ would be significantly worse. Clearly, we need develop heuristics for this problem.

4.2 Recursive Expansion (RE)

The simplest means of efficiently obtaining an A matrix for a given signature tensor S is to make use of the modulo identity $2ab = a+b-a \oplus b$. This is applied recursively to each non-linear term in the weighted polynomial associated with S to yield a decomposition of the form in equation 6. Or explicitly, for each non-zero coefficient in the weighted polynomial $l_\alpha, q_{\alpha,\beta}, c_{\alpha,\beta,\gamma}$, make the following substitutions to the corresponding monomials:

$$x_\alpha \rightarrow x_\alpha, \quad (8)$$

$$2x_\alpha x_\beta \rightarrow x_\alpha + x_\beta - x_\alpha \oplus x_\beta, \quad (9)$$

$$4x_\alpha x_\beta x_\gamma \rightarrow x_\alpha + x_\beta + x_\gamma - x_\alpha \oplus x_\beta - x_\alpha \oplus x_\gamma - x_\beta \oplus x_\gamma + x_\alpha \oplus x_\beta \oplus x_\gamma \quad (10)$$

We call this the *recursive expansion* (RE) algorithm, which has been shown to yield worst-case T counts of $\mathcal{O}(n^3)$. It is straightforward to see this, as gate synthesis matrices resulting from the RE algorithm may include any column with a maximum Hamming weight of 3. There are $\sum_{k=1}^3 \binom{n}{k} = \mathcal{O}(n^3)$ such columns so there can be at most $\mathcal{O}(n^3)$ T gates in the corresponding circuit decomposition.

4.3 Temporary Order Reduction (TOR)

An efficient heuristic that requires at most only $\mathcal{O}(n^2)$ T gates was proposed in [6], which involves breaking down an input CNOT + T circuit into a sequence of control- U operators where each successive operator acts on fewer qubits than the previous, eventually allowing the optimal Reed-Muller decoder to be used when $n \leq n_{RM}$ for a particular control- U . Each target U corresponds to a weighted polynomial with only up to quadratic terms, and the control qubit maps to a particular variable in the weighted polynomial that has been factorized out from all terms in which it appears. The reduction in order from a cubic to a quadratic weighted polynomial allows for an optimal decomposition U' for each target U to be found efficiently using Lempel's matrix factoring algorithm [7]. Additional steps are required to obtain a decomposition of the form of equation 6 from the control- U' sequence, which leads to sub-optimality in the general case but with T count scaling of $\mathcal{O}(n^2)$, a factor n improvement over RE. This algorithm we refer to as the *temporary order reduction* (TOR) algorithm efficiently yields optimal T counts for the special case where every gate in the CNOT + T circuit shares a control.

There are two versions of TOR: with and without feedback. This refers to whether or not extra terms that do not depend on the control qubit of the current iteration are immediately synthesized after calling Lempel's algorithm or if they are fed back into the input signature tensor to be synthesized in the subsequent step.

4.4 Third Order Duplicate and Destroy (TODD)

In this section, we present an algorithm based in principle on Lempel's matrix factoring algorithm [7] that is extended to work for tensors of order 3. This algorithm requires some initial A matrix to be generated by another algorithm such as RE or TOR, then it reduces the number of columns the initial gate synthesis matrix iteratively until exit. In section 5, we present numerical evidence that it is the best efficient solver of the 3-STR problem developed so far. The algorithm works by identifying a subset of columns of an input gate synthesis matrix, A , to which we can add an arbitrary column vector, \mathbf{z} . This allows us to force pairs of columns to be identical, and identical columns can be removed without changing the unitary that it implements up to a Clifford factor. We call this the *Third Order Duplicate and Destroy* (TODD) algorithm because, much like the villainous Victorian barber, it shaves away at the columns of the input A matrix iteratively until the algorithm finishes execution.

We begin by introducing the key mechanism through which we reduce the T count of quantum circuits: by *destroying* pairs of duplicate columns of a gate synthesis matrix, a process which can be done without changing the signature tensor, as shown in the following lemma.

Lemma 4.1. *Let $A \in \mathbb{Z}^{(n,m)}$ be a gate synthesis matrix whose a^{th} and b^{th} columns are duplicates and $A' \in \mathbb{Z}^{(n,m-2)}$ be a gate synthesis matrix formed from by removing the a^{th} and b^{th} columns of A . It follows that $S^{(A)} = S^{(A')}$ for any such A and A' .*

Proof. We start by writing the signature tensor in terms of the elements of A according to equation 7,

$$S_{\alpha,\beta,\gamma}^{(A)} = \sum_{k=1}^m A_{\alpha,k} A_{\beta,k} A_{\gamma,k} \pmod{2}, \quad (11)$$

and separating the terms associated with a, b from the rest of the summation,

$$S_{\alpha,\beta,\gamma}^{(A)} = \sum_{j \in \mathcal{J}} A_{\alpha,j} A_{\beta,j} A_{\gamma,j} + A_{\alpha,a} A_{\beta,a} A_{\gamma,a} + A_{\alpha,b} A_{\beta,b} A_{\gamma,b} \pmod{2}, \quad (12)$$

where $\mathcal{J} = [1, m] \setminus \{a, b\}$. As stated in the lemma, the a^{th} and b^{th} columns of A are duplicates or formally this corresponds to the assertion that

$$A_{i,a} = A_{i,b} \quad \forall i \in [1, n]. \quad (13)$$

Now substitute equation 13 into equation 12,

$$S_{\alpha,\beta,\gamma}^{(A)} = \sum_{j \in \mathcal{J}} A_{\alpha,j} A_{\beta,j} A_{\gamma,j} + A_{\alpha,a} A_{\beta,a} A_{\gamma,a} + A_{\alpha,a} A_{\beta,a} A_{\gamma,a} \pmod{2} \quad (14)$$

$$= \sum_{j \in \mathcal{J}} A_{\alpha,j} A_{\beta,j} A_{\gamma,j} \pmod{2} \quad (15)$$

$$= S_{\alpha,\beta,\gamma}^{(A')}, \quad (16)$$

where the penultimate step follows from modulo 2 addition and the final step follows from the definition of A' and equation 7. \square

Lemma 4.1 gives us a simple means to remove columns from a gate synthesis matrix by destroying pairs of duplicates columns and thereby reducing the T count of a CNOT + T circuit by 2. However, it is often the case that the A matrix does not already contain any duplicate columns. Therefore, we wish perform some transformation on the A matrix, which results in a different gate synthesis matrix A' that: a) has duplicate columns; b) preserves the signature tensor of A ; and c) does not increase the column number of A by more than 1. In the following lemma we introduce a class of transformations that *duplicate* a particular column of an A matrix such that a) and c) are satisfied. We then use lemma 4.3 establish what conditions must be satisfied for the duplication transformation to have property b).

Lemma 4.2. *Let A be an $n \times m$ gate synthesis matrix where all columns are unique and non-zero. Let $A' = A \oplus \mathbf{z}\mathbf{y}^T$ where \mathbf{z} and \mathbf{y} are column vectors on \mathbb{Z}_2 of length n and m , respectively, defined such that $z_i = A_{i,a} \oplus A_{i,b}$ for some $a, b \in [1, m]$ and $y_a \oplus y_b = 1$. It follows that the a^{th} and b^{th} columns of A' are duplicates.*

Proof. We begin the proof by finding expressions for the matrix elements of A' in terms of A , \mathbf{z} and \mathbf{y} ,

$$A'_{i,j} = A_{i,j} \oplus z_i y_j, \quad (17)$$

and substitute the definition of \mathbf{z} ,

$$A'_{i,j} = A_{i,j} \oplus (A_{i,a} \oplus A_{i,b}) y_j. \quad (18)$$

Now we can find the elements of the columns a and b of A' ,

$$A'_{i,a} = A_{i,a} \oplus (A_{i,a} \oplus A_{i,b}) y_a, \quad (19)$$

$$A'_{i,b} = A_{i,b} \oplus (A_{i,a} \oplus A_{i,b}) y_b. \quad (20)$$

We substitute in the condition $y_b = y_a \oplus 1$ into 20,

$$\begin{aligned} A'_{i,b} &= A_{i,b} \oplus (A_{i,a} \oplus A_{i,b})(y_a \oplus 1) \\ &= A_{i,a} \oplus (A_{i,a} \oplus A_{i,b}) y_a \\ &= A'_{i,a}. \end{aligned} \quad (21)$$

\square

Lemma 4.3. *Let A and $A' = A \oplus \mathbf{x}\mathbf{y}^T$ be two gate synthesis matrices where \mathbf{z} , \mathbf{y} are arbitrary column vectors of dimension n and m , respectively. $S(A) = S(A')$ as long as the following conditions hold true:*

$$C1: \quad |\mathbf{y}| = 0 \pmod{2}$$

$$C2: \quad A\mathbf{y} = \mathbf{0}$$

$$C3: \quad \chi(A, \mathbf{z}) \mathbf{y} = \mathbf{0}.$$

where we define $\chi(A, \mathbf{z})$ as follows:

Definition 4.1. χ **Matrix.** Given some gate synthesis matrix, A , and a column vector $\mathbf{z} \in \mathbb{Z}_2^n$ let

$$\chi(A, \mathbf{z}) = \begin{pmatrix} (z_1 \mathbf{r}_2 \wedge \mathbf{r}_3) \oplus (z_2 \mathbf{r}_3 \wedge \mathbf{r}_1) \oplus (z_3 \mathbf{r}_1 \wedge \mathbf{r}_2) \\ (z_1 \mathbf{r}_2 \wedge \mathbf{r}_4) \oplus (z_2 \mathbf{r}_4 \wedge \mathbf{r}_1) \oplus (z_4 \mathbf{r}_1 \wedge \mathbf{r}_2) \\ \vdots \\ (z_{n-2} \mathbf{r}_{n-1} \wedge \mathbf{r}_n) \oplus (z_{n-1} \mathbf{r}_n \wedge \mathbf{r}_{n-2}) \oplus (z_n \mathbf{r}_{n-2} \wedge \mathbf{r}_{n-1}) \end{pmatrix},$$

where \mathbf{r}_i is the i^{th} row of A , and $\mathbf{x} \wedge \mathbf{y}$ is the element-wise product of vectors \mathbf{x} and \mathbf{y} .

Proof. We begin the proof by finding an expression for $S(A')$ using equation 7,

$$S_{\alpha,\beta,\gamma}^{(A')} = \sum_{j=1}^m (A_{\alpha,j} \oplus z_{\alpha} y_j) (A_{\beta,j} \oplus z_{\beta} y_j) (A_{\gamma,j} \oplus z_{\gamma} y_j) \pmod{2}, \quad (22)$$

and expanding the brackets,

$$\begin{aligned} S_{\alpha,\beta,\gamma}^{(A')} = \sum_{j=1}^m & (A_{\alpha,j} A_{\beta,j} A_{\gamma,j} \oplus z_{\alpha} z_{\beta} z_{\gamma} y_j \\ & \oplus z_{\alpha} z_{\beta} A_{\gamma,j} y_j \oplus z_{\beta} z_{\gamma} A_{\alpha,j} y_j \oplus z_{\gamma} z_{\alpha} A_{\beta,j} y_j \\ & \oplus z_{\alpha} A_{\beta,j} A_{\gamma,j} y_j \oplus z_{\beta} A_{\gamma,j} A_{\alpha,j} y_j \oplus z_{\gamma} A_{\alpha,j} A_{\beta,j} y_j) \pmod{2}. \end{aligned} \quad (23)$$

We can see that the first term of equation 23 summed over all j is equal to $S(A)$, by definition. The task is to show that the remaining terms sum to zero under the specified conditions. Next, we use the definitions of $|\mathbf{y}|$, $A\mathbf{y}$ and $\chi(A, \mathbf{z}) \mathbf{z}$ to simplify the result,

$$S_{\alpha,\beta,\gamma}^{(A')} = S_{\alpha,\beta,\gamma}^{(A)} \oplus z_{\alpha} z_{\beta} z_{\gamma} |\mathbf{y}| \oplus z_{\alpha} z_{\beta} [A\mathbf{y}]_{\gamma} \oplus z_{\beta} z_{\gamma} [A\mathbf{y}]_{\alpha} \oplus z_{\gamma} z_{\alpha} [A\mathbf{y}]_{\beta} \oplus [A\chi(A, \mathbf{z}) \mathbf{y}]_{\iota(\alpha,\beta,\gamma)}, \quad (24)$$

where we $\iota(\alpha, \beta, \gamma)$ is the row index of χ whose elements are given by $(z_{\alpha} \mathbf{r}_{\beta} \wedge \mathbf{r}_{\gamma}) \oplus (z_{\beta} \mathbf{r}_{\gamma} \wedge \mathbf{r}_{\alpha}) \oplus (z_{\gamma} \mathbf{r}_{\alpha} \wedge \mathbf{r}_{\beta})$. Applying condition *C1*, the second term is eliminated; by applying condition *C2*, the next three terms are eliminated, and by applying condition *C3*, the final term is eliminated. \square

Now we have shown how to *duplicate* and *destroy* columns of a gate synthesis matrix, we are ready to describe the TODD algorithm, presented as pseudo-code in algorithm 1. Given an input gate synthesis matrix A with signature tensor S , we begin by iterating through all column pairs of A given by indices a, b . We construct the vector $\mathbf{z} = \mathbf{c}_a \oplus \mathbf{c}_b$ where \mathbf{c}_j is the j^{th} column of A , as in lemma 4.2. We check to see if the conditions in lemma 4.3 are satisfied for \mathbf{z} by forming the matrix,

$$\tilde{A} = \begin{pmatrix} A \\ \chi(A, \mathbf{z}) \end{pmatrix}. \quad (25)$$

Any vector, \mathbf{y} , in the null space of \tilde{A} satisfies *C2* and *C3* of lemma 4.3. We scan through the null space basis until we find a \mathbf{y} such that $y_a \oplus y_b = 1$. At this stage we know that we can remove at least one column from A . If $|\mathbf{y}| = 0 \pmod{2}$ then condition *C1* is satisfied and we can perform the duplication transformation from lemma 4.3. Otherwise, we force *C1* to be satisfied by appending a 1 to \mathbf{y} and an all-zero column to A before applying the duplication transformation. Finally, we *duplicate* columns by adding the value of \mathbf{z} to every column j for which $y_j = 1$, then *destroy* the a^{th} and b^{th} columns of A . This reduces the number of columns of A and therefore the T count of U_f . We now start again from the beginning, iterating over columns of the new A matrix. The algorithm terminates if every column pair has been exhausted without success.

Algorithm 1 Third Order Duplicate-then-Destroy (TODD) Algorithm

Input: Gate synthesis matrix $A \in \mathbb{Z}_2^{(n,m)}$.

Output: Gate synthesis matrix $A' \in \mathbb{Z}_2^{(n,m')}$ such that $m' \leq m$ and $S^{(A')} = S^{(A)}$.

- Let $\text{col}_j(A)$ be a function that returns the j^{th} column of A .
- Let $\text{cols}(A)$ be a function that returns the number of columns of A .
- Let $\text{nullspace}(A)$ be a function that returns a matrix whose columns generate the right nullspace of A .
- Let $\text{simplify}(A)$ be a function that removes every pair of identical columns and every all-zero column from A .

procedure TODD

Initialize $A' \leftarrow A$

start:

```
for all  $1 \leq a < b \leq \text{cols}(A')$  do
   $\mathbf{z} \leftarrow \text{col}_a(A') + \text{col}_b(A')$ 
   $\tilde{A} \leftarrow \begin{pmatrix} A' \\ \chi(A', \mathbf{z}) \end{pmatrix}$ 
   $N \leftarrow \text{nullspace}(\tilde{A})$ 
  for all  $1 \leq k \leq \text{cols}(N)$  do
     $\mathbf{y} \leftarrow \text{col}_k(N)$ 
    if  $y_a \oplus y_b = 1$  then
      if  $|\mathbf{y}| = 1 \pmod{2}$  then
         $A' \leftarrow \begin{pmatrix} A' & \mathbf{0} \end{pmatrix}$ 
         $\mathbf{y} \leftarrow \begin{pmatrix} \mathbf{y} \\ 1 \end{pmatrix}$ 
         $A' \leftarrow A' + \mathbf{zy}^T$ 
       $\text{simplify}(A')$ 
    goto start
```



5 Results

Our T gate optimization protocol was implemented in C++ in order to achieve two goals: first benchmark test our different T -Optimise algorithms against one another using random circuits (see figure 3); and second to evaluate the performance of the full Clifford + T protocol against the best known circuit decompositions for useful quantum algorithms, the results of which are listed in table 1. We see from both results that TODD produces quantum circuits with smallest T count.




Figure 3: Circuits generated by the CNOT and T gate were randomly generated for varying number of qubits n then optimized by our implementations of RE, TOR and TODD. The average T -count for each n over many random circuits are shown on the vertical axis. TODD produces circuit decompositions with the smallest T -counts on average but scales the same as the next best algorithm, TOR (Feedback). Both of these algorithms are better than RE by a factor n . The difference between the T -counts for TODD and TOR (Feedback) converge on a constant 5.5 ± 0.7 for large n .

Table 1: T -counts for various universal Clifford + T benchmark circuits as synthesized by the TODD and RE algorithm are shown. T_{original} are the best known results produced by the works cited in the *Circuit* column, and T_{TODD} is the result for *TODD*. n_{original} is the number of qubits of the original circuit and n_{out} is the total number of qubits of the output circuit including ancillas used to implement multiply controlled Toffoli gates as well as Hadamards using path variables [?]. The total execution time in seconds for *TODD* run on an *Intel i7 2.40Gz* processor is given in column *Time*.

Circuit	T_{original}	T_{TODD}	T_{RE}	n	h	Time (s)
hwb6-47-107 [3]	71	55	102	6	43	91.713
hwb6-42-150 [3]	71	46	140	6	43	160.198
nth_prime6_inc_55_667 [3]	400	263	354	6	39	228.025
ham15-109-214 [3]	97	28	65	15	47	27.756
ham15-70 [3]	230	103	148	15	47	100.899
ham15tc1 [3]	1019	258	359	15	50	270.862
#0117 [8]	79	13	63	6	116	118.64
#017F [8]	80	23	59	6	59	20.618
#0001 [8]	40	21	46	6	39	5.2
#001F [8]	43	22	55	6	39	4.382
#0007 [8]	47	13	31	6	64	10.141
#007F [8]	40	19	46	6	27	1.267

6 Discussion

7 Conclusion

8 Acknowledgements

References

- [1] A. Montanaro, “Quantum circuits and low-degree polynomials over \mathbb{F}_2 ,” *Journal of Physics A: Mathematical and Theoretical*, vol. 50, no. 8, p. 084002, 2017. [Online]. Available: <http://stacks.iop.org/1751-8121/50/i=8/a=084002>
- [2] M. J. Bremner, R. Jozsa, and D. J. Shepherd, “Classical simulation of commuting quantum computations implies collapse of the polynomial hierarchy,” *Proceedings of the Royal Society of London A: Mathematical, Physical and Engineering Sciences*, 2010. [Online]. Available: <http://rspa.royalsocietypublishing.org/content/early/2010/08/05/rspa.2010.0301>
- [3] M. Amy and M. Mosca, “T-count optimization and Reed-Muller codes,” *arXiv*, vol. arXiv:1601.07363v1 [quant-ph], 2016.
- [4] M. Amy, D. Maslov, M. Mosca, and M. Roetteler, “A meet-in-the-middle algorithm for fast synthesis of depth-optimal quantum circuits,” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 32, no. 6, pp. 818–830, June 2013.
- [5] M. Amy, D. Maslov, and M. Mosca, “Polynomial-time T-depth optimization of Clifford+T circuits via matroid partitioning,” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 33, no. 10, pp. 1476–1489, Oct 2014.
- [6] E. T. Campbell and M. Howard, “Unified framework for magic state distillation and multiqubit gate synthesis with reduced resource cost,” *Phys. Rev. A*, vol. 95, no. 022316, 2017.
- [7] A. Lempel, “Matrix factorization over $\text{GF}(2)$ and trace-orthogonal bases of $\text{GF}(2)$,” *SIAM J. Comput.*, vol. 4, no. 2, 1975.
- [8] M. Soeken, M. Roetteler, N. Wiebe, and G. D. Micheli, “Logic synthesis for quantum computing,” *arXiv preprint*, vol. arXiv:1706.02721, 2017.